RingCentral AI APIs for Conversational Intelligence

Transcript Report

Source Audio File: sample.2.txt

Cleaned Transcript

Hello everyone, my name is Tom. And today I have with me two of my awesome colleagues, and we can start by getting them introduced. Sameer.

Hey everyone, this is Sameer, and I help in articulating the value provided by our open platform and the plethora of use cases that can be supported. Let's go to the next one. Hi everyone, my name is Will, and I talk about how we use the APIs, how we integrate them, how we deploy them, and go live in our production applications. Awesome, thanks, guys.

Today, the topic is about RingCentral's AI APIs around conversational intelligence. So, let's get started. Question for Sameer: AI is a very highly used acronym over the past four to five years, especially.

How do you understand it in context of conversations? What capabilities does it bring? And how easy is it to use? Very well, Tom, you hit the nail on the head. Yes, AI is a very, very highly used acronym for the last, I would say, five to six years. But when you think about AI in the context of conversations, the first thing that comes to mind is speech-to-text capability or text-to-speech capabilities, right? But I think it's much, much beyond that.

For example, you can do a lot of interaction analysis, sentiment analysis, and emotion analysis with these APIs that we offer. It means if you are in a contact center kind of a use case where a supervisor needs to understand what kind of calls are coming in, right? From all the people around the world, to understand the sentiment level. They can simply have a pointer zero to five to figure out what range the incoming calls fall into, to understand the sentiment of those calls, right? That's one of the biggest, biggest powers that we give to our AI APIs.

Wow, that is so interesting. Next question I have here is for Will: What are the different sets of APIs that we support today? And can you throw some light on the three of the topmost used conversational intelligence API use cases across our portfolio of API products? Yeah, great question.

We basically support APIs in two categories. One, they're audio APIs, and the other set of APIs are based on text. And the three most popular APIs are, you can say, speech-to-text, because even audio gets converted into text, and then the AI engine gives you the transcript of that.

And it can give you a transcript with punctuation, such as commas, so that you can turn it into a report. Another very popular API is the speaker diarization API. So, for example, there are multiple people here in this meeting, and the API will tell you which speaker is speaking at which point in time, who spoke what, essentially.

And then there is another API that does speaker identification. This is similar to the speaker diarization API. It detects who's speaking.

And if the AI model has been trained by the speaker's voice, it can tell you who the person is if you provide a label, such as the name of the person. And then it will tell you, for example, 'Will is speaking at this time.' Otherwise, it will just say 'Person A,' 'Person B,' 'Person C.' That's great, Will.

Sameer, Will, thank you so much again for explaining the whys and the hows of RingCentral's open platform and specifically our intelligent APIs. This program is in beta right now. We'll soon be working toward releasing it generally available to the public. So please stay tuned, and we look forward to getting some feedback from you.